

Project 2

CE/CZ4042: Neural Networks and Deep Learning

Deadline: 15 Nov 2019

- You need to complete both parts A and B of the project and submit a project report and source codes online via NTULearn before the deadline.
- Data files for both parts are found in Project 2 folder under Assignments on NTULearn.
- The project is to be done in a group of not more than two students. Both members of the group should submit the project report and codes to NTU Learn, using their individual accounts. The cover page of the report should contain the names of both members.
- The project report would contain sections on (i) introduction, (ii) methods, (iii) experiments and results, and (iv) conclusions. The report should be submitted in the
 - lastname_firstname_P2_report.pdf (report in pdf format); and
 - lastname_firstname_P2_codes.zip (all the source codes in a zip file)
- The assessment will be based on both the project report and the correctness of the codes submitted. You are encouraged to use TensorFlow to complete the project but using other libraries is allowed. Late submissions will be penalized: 5% for each day up to three days.
- This project requires the use of GPU enabled machines. You can access these machines in Software Lab 3 located at N4-B1c-14 on
 - Mondays 12.30pm-5.30pm (Weekly),
 - Tuesdays 8.30am-2.30pm (Odd weeks),
 - Wednesdays 8.30am-12:30pm (Even weeks),
 - Thursdays 10.30am-12.30pm/2.30pm-5pm (Odd weeks), and
 - Fridays 12.30pm-5pm (Even weeks).Please contact the lab manager Mr. Goh Tong Hai (THGoh@ntu.edu.sg) for remote access. Note that only 10PCs are allocated for remote access.
- TAs Mr. Kelvin Chan (CHAN0899@e.ntu.edu.sg) and Mr. Wenwei Zhang (WENWEI001@e.ntu.edu.sg) are in charge of the course project. Please see them at MICL (N4-B1c-17) or PDCL (N4-B2a-02) during their office hours: Friday 3pm-4pm, in case you face issues.

Part A: Object Recognition

The project uses a sample of the CIFAR-10 dataset:

<https://www.cs.toronto.edu/~kriz/cifar.html>

The dataset contains RGB colour images of size 32 X 32 and their labels from 0 to 9. You will be using the batch-1 of the dataset, which contains 10,000 training samples. Testing is done on 2,000 test samples. The training data and testing data are provided in files 'data_batch_1' and 'test_batch_trim' files, respectively. Sample code is given in file start_project_2a.py

Design a convolutional neural network consisting of:

- An Input layer of 3x32x32 dimensions
- A convolution layer C_1 of 50 filters of window size 9x9, VALID padding, and ReLU neurons. A max pooling layer S_1 with a pooling window of size 2x2, with stride = 2 and padding = 'VALID'.
- A convolution layer C_2 of 60 filters of window size 5x5, VALID padding, and ReLU neurons. A max pooling layer S_2 with a pooling window of size 2x2, with stride = 2 and padding = 'VALID'.
- A fully connected layer F_3 of size 300.
- A softmax layer F_4 of size 10.

1. Train the network by using mini-batch gradient descent learning. Set batch size =128, and learning rate $\alpha = 0.001$. Images should be scaled.
 - a. Plot the training cost and the test accuracy against learning epochs.
 - b. For any two test patterns, plot the feature maps at both convolution layers (C_1 and C_2) and pooling layers (S_1 and S_2) along with the test patterns.

(18 marks)

2. Using a grid search, find the optimal numbers of feature maps for part (1) at the convolution layers. Use the test accuracy to determine the optimal number of feature maps.

(10 marks)

3. Using the optimal number of filters found in part (2), train the network by:
 - a. Adding the momentum term with momentum $\gamma = 0.1$.
 - b. Using RMSProp algorithm for learning
 - c. Using Adam optimizer for learning
 - d. Adding dropout to the layers

Plot the training costs and test accuracies against epochs for each case.

(12 marks)

4. Compare the accuracies of all the models from parts (1) - (3) and discuss their performances.

(10 marks)

Part B: Text classification

The dataset used in this project contains the first paragraphs collected from Wikipage entries and the corresponding labels about their category. You will implement CNN and RNN layers at the word and character levels for the classification of texts in the paragraphs. The output layer of the networks is a **softmax layer**.

The training and test datasets will be read from 'train_medium.csv' and 'test_medium.csv' files. The **training dataset** contains **5600 entries** and **test dataset** contains **700 entries**. The **label** of an entry is **one of the 15 categories** such as people, company, schools, etc.

The input data is in text, which should be converted to character/word IDs to feed to the networks by using 'tf.contrib.learn.preprocessing'. **Restrict the maximum length of the characters/word inputs to 100**. Use the **Adam optimizer for training** with a **batch size = 128** and **learning rate = 0.01**. Assume there are **256 different characters**.

1. Design a Character CNN Classifier that receives character ids and classifies the input. The CNN has two convolution and pooling layers:
 - A convolution layer C_1 of 10 filters of window size 20x256, VALID padding, and ReLU neurons. A max pooling layer S_1 with a pooling window of size 4x4, with stride = 2, and padding = 'SAME'.
 - A convolution layer C_2 of 10 filters of window size 20x1, VALID padding, and ReLU neurons. A max pooling layer S_2 with a pooling window of size 4x4, with stride = 2 and padding = 'SAME'.

Plot the entropy cost on the training data and the accuracy on the testing data against training epochs.

(9 marks)

2. Design a Word CNN Classifier that receives word ids and classifies the input. Pass the inputs through an embedding layer of size 20 before feeding to the CNN. The CNN has two convolution and pooling layers with the following characteristics:
 - A convolution layer C_1 of 10 filters of window size 20x20, VALID padding, and ReLU neurons. A max pooling layer S_1 with a pooling window of size 4x4, with stride = 2 and padding = 'SAME'.
 - A convolution layer C_2 of 10 filters of window size 20x1, , VALID padding, and ReLU neurons. A max pooling layer S_2 with a pooling window of size 4x4, with stride = 2 and padding = 'SAME'.

Plot the entropy cost on training data and the accuracy on testing data against training epochs.

(9 marks)

3. Design a Character RNN Classifier that receives character ids and classify the input. The RNN is GRU layer and has a hidden-layer size of 20.

Plot the entropy cost on training data and the accuracy on testing data against training epochs.

(9 marks)

4. Design a word RNN classifier that receives word ids and classify the input. The RNN is GRU layer and has a hidden-layer size of 20. Pass the inputs through an embedding layer of size 20 before feeding to the RNN.

Plot the entropy on training data and the accuracy on testing data versus training epochs.

(9 marks)

5. Compare the test accuracies and the running times of the networks implemented in parts (1) – (4).

Experiment with adding dropout to the layers of networks in parts (1) – (4), and report the test accuracies. Compare and comment on the accuracies of the networks with/without dropout.

(8 marks)

6. For RNN networks implemented in (3) and (4), perform the following experiments with the aim of improving performances, compare the accuracies and report your findings:

- Replace the GRU layer with (i) a vanilla RNN layer and (ii) a LSTM layer
- Increase the number of RNN layers to 2 layers
- Add gradient clipping to RNN training with clipping threshold = 2.

(8 marks)

Hint: Sample code is given in file start_project_2b1.py and start_project_2b2.py